

CSci551 Fall 2014 Project A

Project A Assigned: 2014-09-13

Practical Project A Due: **noon, Wed. 2014-10-01** (not a class day).

Research Project A due: **noon, Wed. 2013-10-01** (not a class day).

You are welcome to discuss your project with other students at the conceptual level, and to use external libraries *after getting permission*. Any cases of plagiarism will result in an F for the entire course. **If you have any questions about policies about academic integrity, please talk to the professor.**

Changes: 2014-09-12: none yet. *But*, you should *expect* that there will be changes and clarifications to this project, and code accordingly. (Start early, but be able to change things slightly as the work goes along.)

On change you may expect is that we will probably downselect to either VirtualBox or VMWare, but not both. To get started you may use either.

1 Overview

There are two *paths* for projects in CSci551 this year. The *Practical Path* project involves implementing a network-centric system. Everyone following this path solves the same problem, but of course in slightly different ways.

In the *Research Path*, students do new research working under the guidance of a mentor (usually a senior PhD student).

PhD students are required to do Research Path projects. Other students may choose either path.

In both cases the project is divided into three parts, Projects A, B and C. Each builds on the prior.

1.1 The Practical Path Project

The purpose of the CSci551 project is to get some hands on experience building a substantial distributed application running on a multi-process computing infrastructure, then to use it to study networking. It also has a secondary purpose of implementing a program using network programming (sockets) and processes (fork, in stage 1) and Unix development tools (make).

You may reuse algorithms from textbooks. You may possibly reuse functions from libraries, but if you're using anything other than the C or C++ standard library (libc, STL), or libraries mentioned on the TA's web page, you *must* check with the TA and the professor and identify it in your write-up. You need to check allowed libraries in Project Information on the class moodle. Otherwise, each student is expected to write *all* of his or her code independently. All programs will be subject to automated checking for similarity.

Please note: you should expect to spend a significant amount of time on the class projects this semester, probably at least 20 hours or more when they're all put together (or that much

per project if you're new to C or C++ and network programming). Please plan accordingly. If you leave all the work until the weekend before it is due, you are unlikely to have a successful outcome. That said, Project A is a “warm-up” project. Projects B and C will be much larger.

For the course, you will do three separate but related projects. First (Project A), you need to demonstrate that you can read the config file, do basic process control, and submit a complete assignment. (Project A doesn't really evaluate anything advanced, but it should confirm that everyone is on the same page and is comfortable with our software environment.) Project A has an early due date. Project A should be relatively short for students used to working on Linux; if not, it should help get you up to speed.

In Project B you will use this facility to implement *Minitor*, a simplified version of a TOR-like Onion Router as described in the paper by Dingledine et al. (see [Dingledine04a] in the class syllabus). It will probably be due just after the midterm. Project B will be *much* larger than Project A; you should plan your time accordingly. Project B will be assigned later and may overlap partially with Project A.

Project C will be assigned later. It will be smaller than Project B but bigger than Project A. It will build on Project B. We will offer a the TA's implementation of Project B for students who wish to use it instead of their own Project B implementation, but we do not promise to help you understand it. Project C will involve extending your Project B Minitor implementation with additional features, and perhaps measuring it or trying to de-anonymize it. It will probably be due the last week of class.

1.2 Research Path Projects

This year we have an alternate *research path* for CSci551 projects. Students will conduct original research on topics related to the course material, helping to further our understanding or develop new approaches for open problems—at least, as much as one can within a semester! For these projects, you can use the language and libraries of your choice, subject to approval from your mentor (more on mentors below). As usual, you must properly credit and cite any code, text, or approaches that you borrow from others. Using and extending the research of others is a crucial part of research, but requires proper attribution.

All PhD students are *required* to do research path projects. MS students have the *option* of doing it, although MS students can cancel research path and pick up the practical project as an alternative after discussing this choice with the professor.

We will have a poster session at the end of the semester, where students can present their research path results. We encourage all students—not just research path students—to attend.

Research Projects in the Friday section are *individual* projects (not groups), but there may be some linked projects where two people work on different parts of the same problem. (Linked projects will share ideas and possibly some data or code or results, but each student will have a specific part they are responsible for and each student must do their own writeup.)

Research Projects in the Tuesday/Thursday section will allow two-person groups projects or individual projects.

2 Practical Project A: Getting Started

2.1 Virtual Machines and your Development Environment

To provide a consistent environment that you can develop your class code on, we will provide you a virtual machine image with Fedora-20 (Linux 3.15.10-201) installed upon it. You can download it from “Project Information” on the class moodle (at <https://moodle.ant.isi.edu/mod/page/view.php?id=241>). Warning: it is 2.1 GB in size to download, and 5.1 GB when unpacked. You are welcome to develop your code on any platform (if you want), but your code *must* build and run on this platform, as it is the only one we will use for evaluation.

There are virtual machines available that can be run with VirtualBox (available at no cost for Windows, Linux, and Mac from <https://www.virtualbox.org>). We believe they also run on VMWare Player (available at no cost for Windows and Linux from <http://www.vmware.com/products/player/>). With some work you could also probably run our VM images under other VM hypervisors (kvm under Linux, or Hyper-V on Windows), although we don’t directly support things other than VirtualBox. (In the past we had some trouble with subtly different operation under different hypervisors.) Please see the TA’s Project Information page for more information about VMs.

Although we recommend VirtualBox, our current instructions are for VMWare Player. We expect to update these shortly, and you may use either in the interim.

Computers on the first floor of SAL (SAL 127 and also the group study area near front desk), have VMWare player pre-installed and you may run it there. (We are evaluating if they can also support VirtualBox, or if not we may add support for VMWare.) Since the lab supports only VMWare and we test under VirtualBox, please e-mail the TA and professor if you plan to use the lab instead of your own computer, so we can check for incompatibilities. Please note that when you submit your code, you will submit *only* your code, not the entire virtual machine. So your code must run without any OS changes.

Accounts in the VM: In the virtual machine, we have created a user account `csci551`, with password `csci551`. We recommend you do all your development as this user (not as root, since it’s very bad security practice to run as root). We also recommend you change the password after you log in. Some parts of the project will require root access. You may use the `sudo` program to get root access using this user password. (We are not giving out the root password, but of course, with full sudo access you can certainly change that if you insist, although you should not need to do so.)

You will need to get files in to your VM, and later you will need to get them out (say, to submit your assignment). The VM has networking turned on, so you should be able to scp and ftp in and out, as well as contact the class moodle from within the VM to submit your project. You may also find it helpful to use a “share folder” to move files between the host computer and the VM. Other ways are possible; Google can help if you want.

The VM starts up with a very basic interface. After logging in, right-click anywhere on the desktop to open a menu containing a list of applications. You will need to run multiple processes on the VM for the project, so you will probably want multiple windows. There are several ways to get multiple windows. You can start separate terminals on the host computer and ssh in to the VM from each. Alternatively, you can run multiple windows inside the VM

using `xterm`.

Other tools: You are welcome to install other development tools if you want. A Google search on “installing Fedora software” will suggest how to do so. Please remember, though, any libraries must be pre-approved, as must any development tools that are required to build your program. If your program requires any libraries or non-pre-installed development tools, you must document that in your program’s README, described below.

Please see Section 7 for other requirements about your project source code.

2.2 Stage 1: Starting Processes

In the first stage you will experiment with process forking and socket communication, and network plumbing, and you will do this all in a virtual machine.

You will create a Onion Proxy program that is responsible for forking an Onion Router. The Proxy will read a configuration file to tell it how many routers to fork. The Proxy will hook in to the networking system through a *tunnel* network interface that lets its scoop up network traffic.

In a later stage the two will send traffic.

The router will wait for traffic on a UDP socket, and generate return traffic the same way.

We will give you the exact commands to hook up network traffic to your proxy, and a framework for handling packet processing in the proxy. You will need to implement packet processing at the router.

Your program will begin by reading a configuration file.

The configuration file is in a simple format with comments and non-comment lines. Any line that starts with a hash character (“#”) is a comment and should be skipped. The first non-comment line will indicate what stage is being run. Initially, this will be **stage 1**. The second non-comment line will contain the text `num_routers N`, where N will be replaced by an integer indicating how many routers to create. For stage 1 (and, in fact, for all stages in Project A), *N* will *always* be 1, but we will change that in later stages.

The Proxy will create a log file “stage1.proxy.out” where it will write information we will use to evaluate your assignment. We will ignore anything printed to standard output or standard error; we *only* evaluate the log file.

The Proxy should create a dynamic (operating-system assigned) UDP port. (See the *getsockname()* man page or the Steven’s book on socket programming in the class syllabus for details about how to do dynamic port allocation if you’re not familiar with it.) After starting the Routers the Proxy will wait on this port for return traffic. The Proxy should print the dynamic port to its log file the line “proxy port: PORT” where PORT is the number of the dynamic port.

After reading the configuration file, your Proxy should `fork()` to create the router. The code for both your proxy and router need to be in the same program, so that after forking, the router should run a router subroutine and the proxy a proxy subroutine. You should not execute another program (you don’t need to, and it would just complicate things).

The Proxy should pass its dynamically allocated UDP port process to the child through a global variable. A copy of this global variable will be available in the child process so it

knows how to contact the Proxy after starting its Router duties. You may also pass the stage number and the router index numbers through global variables, if you want.

The Proxy will then wait for the Router to send it an “I’m up” message on its UDP port. You may select whatever format you want for this message.

When the Router starts, it should create a log file called “stage1.router1.out”. It should then get its process id (see the *getpid()* man page) and write the line “router: X, pid: N, port: M” to this file. It should then send an “I’m up” message (Just send router’s PID) to the Proxy’s UDP socket.

After forking, the Proxy will wait passively on its UDP socket for packets from routers. When a router sends “I’m up” message to the Proxy, the Proxy should first print the line “router: X, pid: N, port: M” to its log file (the Proxy’s log file).

All code must run in the same way. Your program should take the configure file as a parameter, so the command to run your code should be like: `sudo ./proja config_file` (This requirement also applies to stage 2).

Sample input: Here is a sample configuration input for stage 1.

```
#This is a comment line, should be ignored
stage 1
num_routers 1
```

Sample output: Here is a sample output for stage 1.

In stage1.proxy.out:

```
proxy port: 44051
router: 1, pid: 3104, port: 43031
```

In stage1.router1.out:

```
router: 1, pid: 3104, port: 43031
```

Writeup: In addition to turning in your code, you need to write up what you learned. Put this information in a text file called `README.stage1.txt`, and these details (include the letters, too, so we know what part of your README answers which of our questions):

- a) **Reused Code:** Did you use code from anywhere for your project? If not, say so. If so, say what functions and where they’re from. (Also identify this with a comment in the source code.) If you use the class timer code, you must say so here and describe any changes you had to make to it.
- b) **Complete:** Did you complete this stage? If not, what works and what doesn’t?
- c) **Portable:** Will your code always work if the proxy and the router were on different computers with different CPU architectures, like an IBM PowerPC and an AMD x86-64? If so, why (what specifically did you do to support that case)? If not, why not (what problem would occur if run between different types of computers)? (Note that in your Project A, they are always on the same computer architecture because you’re

just forking another process, not connecting to another computer. This question is about the hypothetical case of if they were on different computers with different CPU architectures.)

Other comments: Even at this stage your program needs to be modular, using multiple files for logically different parts of the source code. One strategy might be to put the proxy and the router each in different files, or you can choose other organizations.

2.3 Stage 2: Sending Traffic

In this stage, we add traffic and hook your proxy and router into the network.

Your proxy will hook into the networking system. To do this, we first need to create a tunnel and configure it. To do so, run the following commands in a terminal:

```
sudo ip tuntap add dev tun1 mode tun
sudo ifconfig tun1 10.5.51.2/24 up
```

The first command creates a new tunnel interface named `tun1`. The second command sets the IP address of the other end of the tunnel to 10.5.51.2, and brings it up. After doing these steps, all the traffic sent to the 10.5.51/24 address block will go into `tun1`. So if you ping address 10.5.51.x from a new terminal, all the ICMP packets will go into `tun1`.

After you create and configure the tunnel, you need to write code to receive tunneled packets. To do this, your program will open the tunnel interface and then busy loop reading packets from the tunnel device. Each packet will be returned in a separate read call.

The code to handle this process is fairly standard; we will provide a framework on moodle that we recommend you use. (Projects/sample_tunnel.c) (Our framework is derived from the tutorial at <http://backreference.org/2010/03/26/tuntap-interface-tutorial/>, if you want more detailed information about how to open a tunnel interface and how to read data from and write data to it.)

Our framework only supports setting up, reading from the tunnel. You must extend this framework to read *both* from the tunnel *and* to get packets over UDP from the router. To do handle both, you will need to replace the busy loop reading the tunnel with a `select()` statement that handles both the tunnel interface and the UDP socket. You hopefully had experience with `select` in EE450; if not, the Stevens textbook on network programming (optional for the class) and a number of web-based tutorials can get you started.

After you get an ICMP ECHO packet from tunnel interface, your proxy will need to send it to the router. To do this, you must send the packet via the UDP socket you created in stage 1.

Your router should receive the ping packet and generate an ICMP echo reply, and send it back to proxy via the UDP socket. You will need to use the standard ICMP packet format for echo reply as defined in RFC-792 (a copy is at <http://tools.ietf.org/html/rfc792>).

Finally, your proxy should write the ICMP echo reply packet to the tunnel interface.

The proxy must log each ICMP packet it forwards (both to the user and to the router) to its log file, with the format “ICMP from tunnel, src: 10.5.51.2, dst: 10.5.51.3, type: 8”, for traffic from the tunnel. Or if from a router, it should say “...from port: N...”

After you did all of the above steps, when you ping any address in 10.5.51/24, you should be able to see replies. **Note:** your program must be run as root user or superuser privileges (for example, by `sudo ./proja config_file`), since Linux requires root privileges to manipulate tunnel interface.

Sample input: Here is a sample configuration input for stage 2.

```
#This is a comment line, should be ignored
stage 2
num_routers 1
```

Sample output: After running your program and then running `ping -c 4 10.5.51.3` in another window:

In stage2.proxy.out:

```
proxy port: 44051
router: 1, pid: 3104, port: 43031
ICMP from tunnel, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 43031, src: 10.5.51.3, dst: 10.5.51.2, type: 8
ICMP from tunnel, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 43031, src: 10.5.51.3, dst: 10.5.51.2, type: 8
ICMP from tunnel, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 43031, src: 10.5.51.3, dst: 10.5.51.2, type: 8
ICMP from tunnel, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 43031, src: 10.5.51.3, dst: 10.5.51.2, type: 8
```

In stage2.router1.out:

```
router: 1, pid: 3104, port: 43031
ICMP from port: 44051, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 44051, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 44051, src: 10.5.51.2, dst: 10.5.51.3, type: 8
ICMP from port: 44051, src: 10.5.51.2, dst: 10.5.51.3, type: 8
```

Writeup: In addition to turning in your code, you need to write up what you learned. Put this information in a text file called `README.stage2.txt`, and these details (include the letters, too, so we know what part of your README answers which of our questions):

- a) **Reused Code:** Did you use code from anywhere for your project? If not, say so. If so, say what functions and where they're from. (Also identify this with a comment in the source code.) If you use the class timer code, you must say so here and describe any changes you had to make to it.
- b) **Complete:** Did you complete this stage? If not, what works and what doesn't?

- c) **Portable:** Will your code always work if the proxy and the router were on different computers with different CPU architectures, like an IBM PowerPC and an AMD x86-64? If so, why (what specifically did you do to support that case)? If not, why not (what problem would occur if run between different types of computers)? (Note that in your Project A, they are always on the same computer architecture because you're just forking another process, not connecting to another computer. This question is about the hypothetical case of if they were on different computers with different CPU architectures.)

3 Research Project A: Getting Started

For Research Project A, students are expected to identify a research topic. We provide you with a list of candidate topics in a Google Doc, posted by 2014-09-16 on the TA's web page. Students may also propose a topic of their choice. In that case, you need to contact the TAs well before the proposal deadline, as you will need to find a mentor willing to work with you on the project. As such, the list of candidate topics we provide may be the easier path.

Each topic will have a mentor. Interested students should contact that mentor by e-mail (cc'ing the professor and TA) and discuss what is involved, what tools will be used, and what the research questions are.

The main goal of this interaction is to make sure that both you and your mentor agree that you understand what you want to do, and that the research project is a good fit for your background.

The result of Research Project A should be a short (one page!) project proposal, with the following numbered sections:

1. what research you plan to do (briefly, a paragraph or so).
2. why does it matter?
 - (a) why is it interesting research (to the field)?
 - (b) why is it interesting for you (what will you learn)?
 - (c) what is relationship to other work (briefly, a few sentences or so, ideally with references to prior work or other projects)
3. what specific *deliverables* you will have for Project B. These must include:
 - (a) meeting once a week with your mentor.
 - (b) written weekly notes about your progress (suggested format: a Google docs document that you can share). This should list: (i) what I accomplished last week, (ii) what questions I have and what problems I need help with, (iii) what I plan to do next week, (iv) comments from the mentor. Each week, please prepare these notes and send them to your mentor the night before your meeting, then update the notes after the meeting (with answers to your questions and any revisions to your plan).

- (c) what end results you will have. Probably some specific code, or measurement. Probably including some specific experiment or analysis.
- (d) That you will do a Project B report if your Research Project B is approved. The Research Project B report is about 2-6 pages summarizing where you are as of the Project B deadline and what (if anything) you plan for Research Project C (tentatively: noon Wed. 2013-10-29).
- (e) If your Research Project C is approved, that you will do a Project C report and an end-of-semester poster. A Research Project C report is about 3-8 pages summarizing where you are at the end of the semester (tentatively: noon Wed. 2013-12-03). Research Project C will also have a required poster. The poster session will be outside the class time, so attendance is optional. (We expect to provide support for poster printing; details will follow.)

The TA and professor will need to *approve* all Research Project A proposals before you move on to Research Project B. While we want to support the research track, we may decline proposals if we think there's not a good fit, or if we have too many people doing the same topic, or for other reasons.

If your Research Project A proposal is approved, this proposal will serve as the requirements for your Research Project B. (We will try to comment on your Project A proposal by the end of the 1st week of October.) Further, it will count towards part of your Research Project B grade (and will NOT be graded if you do not do Research Project B).

If we all (you, your mentor, the TA, and the professor) agree your Research Project B was successful then you have the option of doing Research Project C. Exactly what defines "successful" is ultimately up to the TA and professor. Potentially it is: did you meet the objectives in your proposal? But, because research isn't always obvious up front, partly it's a judgment call. It's *your job* to make the case that your research is successful in your Research Project B report. Note that it will take us some time to evaluate Research Project B reports, so you may have to decide which track for Project C with only partial information. If we do not approve your proceeding to Research Project C, you may be required to switch to Practical Project C after the assignment goes out. Since PhD students are required to do Research Project C, we suggest you make sure your project is strong enough to proceed. Please talk to the professor or TA if you are unsure which way to go. (We will try to confirm go-ahead for Project C within a week of Project B completion.)

Your Research Project A proposal, plus feedback you get in Project B, will be the requirements for your Research Project C.

Two observations about the Research Project track. We expect it will be more work than the Practical Project Track. You have to do both Research Project A and Practical Project A (although that's small). And in general, research is more work than just implementing a practical project. Further, for students in the Tuesday/Thursday section who choose to work with a partner, the group will be expected to jointly conduct more work than an individual. However, for students interested in research, the research track should be more interesting as well, and you can expect to learn much more about the research process. (For similar classes, interested students with very successful projects have gone on to turn them into papers, or code accepted in open source projects, although getting to these stages usually requires at least another semester of work.)

Second, the research track is somewhat more risky. Research fails. Sometimes assumptions are wrong. Although we will identify research topics we think are viable, there are both known unknowns and unknown unknowns. For this reason we provide a path to transition from Research to Practical Path for non-PhD students at each stage of the project. If you are on the Research Project Track, you have until the Project B deadline to either do Research Project B or Practical Project B; if you switch to Practical Project B you are committed to doing Practical Project C.

3.1 Submitting your Research Project A Proposal

Your Research Project A proposal is a one-page document that **must** be in plaintext or PDF format. See Section 3 for what needs to be in the document.

Submissions that are in .doc, .docx, .odf, or any other format **will not** be accepted. Only plaintext or PDF formats will be accepted. There are many ways to export to PDF: see PDFCreator from <http://www.pdfcreator.org/pdfcreator> for one way.

To submit your Research Project A proposal, upload your proposal to the class moodle with the filename format if it is a PDF file:

LASTname_FIRSTname_rproja.pdf

If it is a plaintext file, use the following format:

LASTname_FIRSTname_rproja.txt

Warning: when you upload to the moodle, please be careful in that you must both do “Upload a file” *and* do “Save changes”! When you are done you should get a message “File uploaded successfully”, and you should see a list with your file there and an option to “update this file”. If you do *not* see “file uploaded successfully”, then you have *not* completed uploading!

4 Practical Project B: Minitor

Project B will be to implement Minitor, a simplified version of a TOR-like Onion Routing system as described in the paper by Dingledine et al. (see [Dingledine04a] in the class syllabus). You may wish to review that paper to prepare for the project. (You should not need to read that paper for Project A.)

5 Research Project B

Research Project B will involve conducting research throughout this portion of the semester, meeting regularly with your mentor, and submitting writeups of your weekly progress and a longer writeup of all your progress and plan for this portion of the project.

6 Practical Project C: Extending Minitor

7 Submitting and Evaluating Practical Projects

To submit each part the assignment, put *all* the files needed (Makefile, README, all source files, and source to any libraries) in a gzip'ed tar file and upload it to the class moodle with the filename `proja.tar.gz`. *Warning:* when you upload to the moodle, please be careful in that you must both do “Upload a file” *and* do “Save changes”! When you are done you should get a message “File uploaded successfully”, and you should see a list with your file there and an option to “update this file”. If you do *not* see “file uploaded successfully”, then you have *not* completed uploading!

We *strongly* recommend that you confirm that you have included everything needed to build your project by extracting your tarfile in a different directory and building it yourself. (It's easy to miss something if you don't check.)

It is a project requirement that your project come with a Makefile and build with just the make command. To evaluate your project we will type the following command:

```
% make
```

Structure the Makefile such that it creates an executable called `proja`. (Note: *not* `a.out`.) For more information please read the `make(1)` man page. (The Linux we use supports GNU `make`; you may use GNU `make` extensions if you wish.)

We will then run your program using a test configuration file. You can assume that the topology description will be syntactically correct. After running the program, we will grade your project based on the output files created by your program's processes.

It is a project requirement that your implementation be somewhat modular. This means that you should follow good programming practices—keep functions relatively short, use descriptive variable names. You must use at least one header file, and multiple files for different parts of your program code. (The whole project should be broken up into *at least* two C/C++ files. If you have a good file hierarchy in mind you can break up into more files but the divisions should be logical and not just spreading functions into many files.) Indicate in a comment at the front of each file what functions that file contains.

We will consider external libraries, but **it is a project requirement that all external libraries must be explicitly approved by the professor.** Any libraries in the default VM image are suitable (including `libc` and `STL`). A list of approved libraries will be on the TA's webpage on the moodle.

Computer languages other than C or C++ will be considered but *must be approved ahead of time*; please contact the professor and TA if you have an alternative preference. The deadline for approving new computer languages is *one week* before the project due date, so get requests in early. The language must support sockets and process creation. (Please ask *before* you start, we don't want you waste your work.)

Although we provide a complete sample input file and output, final evaluation of your program will include other input sources. We therefore advise you to test your program with a range of inputs.

Although the exact output from your program may be different from the sample output we provide (due to events happening in different orders), your output should match ours in format.

It is a project requirement that your code build without warnings (when compiled with `-Wall`). We will be responsible for making sure there are no warnings in class-provided code (any warnings in class-provided code are our bugs and will not count against you).

8 Hints

The structure of the project is designed to help you by breaking it up into smaller chunks (compared to the size of the whole project). We strongly encourage you to follow this in your implementation, and do the stages in order, testing them as you go. In the past, some students have tried to read and implement the whole assignment all at once, almost always resulting in an unhappy result.

8.1 Common pitfalls

Please do not hardcode any directory path in your code! If you hardcode something like `//home/csci551/...` in your code to access something in your home directory and the grader cannot access these directories during grading, your code will not work (and this will be your fault)! If your code does not work, you get no credit! Instead, assume paths are given external to your program, and that you read and write files in the current directory (wherever that is).

8.2 Doing multiple things at once

Later stages of the project may require you to handle both timers and I/O at the same time. (Stage 1 is not complex enough to require timers.) One approach would be to use threads, but most operating systems and many network applications don't actually use threads because thread overhead can be quite large (not context switch cost, but more often memory cost—most threads take at least 8–24KB of memory, and on a machine with 1000s of active connections that adds up, and always in debugging time, in that you have to deal with synchronization and locking). Instead of threads, we strongly encourage you to use timers and event driven programming with a single thread of control. (See the talk “Why Threads Are A Bad Idea (for most purposes)” by John Ousterhout, <http://www.stanford.edu/~ouster/cgi-bin/papers/threads.pdf> for a more careful explanation of why.)

Creating a timer library from scratch is interesting, but non-trivial. We will provide a timer library that makes it easy to schedule timers in a single-threaded process. You may download this code from the TA web page. There is *no* requirement to use this code, but you may if you want. If you want to use it, download it from the class web page. There is no external documentation, but please read the comments in the `timers.hh` and look at `test-app.cc` as an example. If you do use the code, you must add it to your Makefile and you must document how you used it in your README.

8.3 Other sources of help

You should see the Unix man pages for details about socket APIs, fork, and Makefiles. Try `man foo` where foo is a function or program.

The TA can provide *some* help about Unix APIs and functionality, but it is not his job to read the manual page for you, nor is it his job to teach how to log in and use vi or emacs.

You may wish to get the book *Unix Network Programming*, Volume 1, by W. Richard Stevens, as a reference for how to use sockets and fork (it's a great book). We will *not* cover this material in class.

The README file should not just be a few sentences. You need to take some time to describe what you did and especially anything you didn't do. (Expect the grader to take off more points for things they have to figure out are broken than for known deficiencies that you document.)